

IoT S&S: Real-Time

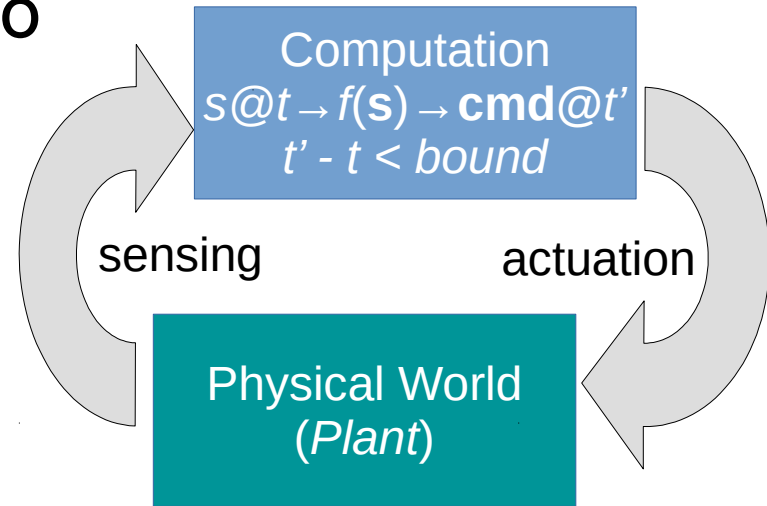
Gabe Parmer



Real-Time/Cyber-Physical Systems

- Embedded systems + **deadlines**

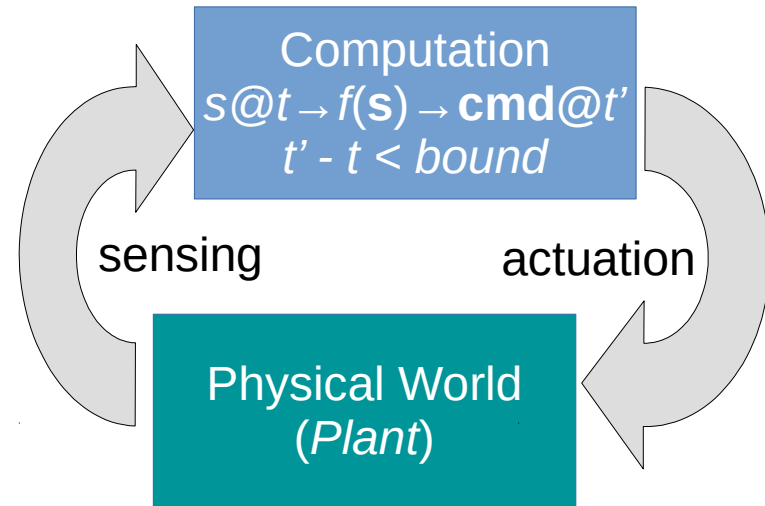
- The computation must adhere to the timing constraints of the physical world
- Send command to actuator within a **bounded time** from receiving sensor information
- We don't write code thinking about time; *how do engineers program real-time systems???*



Real-Time/Cyber-Physical Systems

Complexities

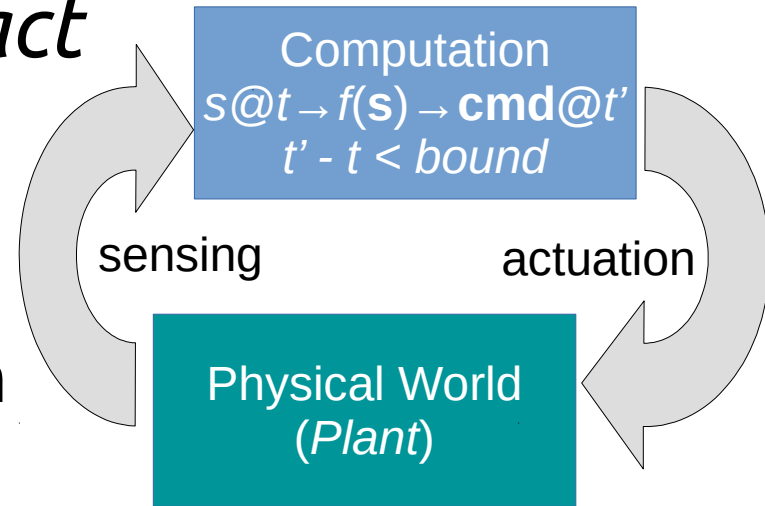
- N sensors, M actuators (s&a)
- Each with a *periodicity* of freshness (sensors) / readiness (acts.)
 - action every t
- Each requires
 - computation every t
 - communication between $s \rightarrow a$



Real-Time/Cyber-Physical Systems

Embedded systems + **deadlines**

- How do we *think about/abstract time* in our computation?
- Task model:
 - Observation – s&a computation is *periodic*
 - Some computation every $t \dots$
how much computation?




Periodic Task Model

Task (τ_i) – abstraction of computation, each has a


- **Period** (p_i) – period, span of repetitive exec.
A task *activates* (wakes up) every p_i time units
- **WCET** (e_i) – worst-case execution time
Every activation, task τ_i executes for max e_i
- **Deadline** of τ_i is p_i after activation

System is $T = \{\tau_0, \tau_1, \dots, \tau_N\}$

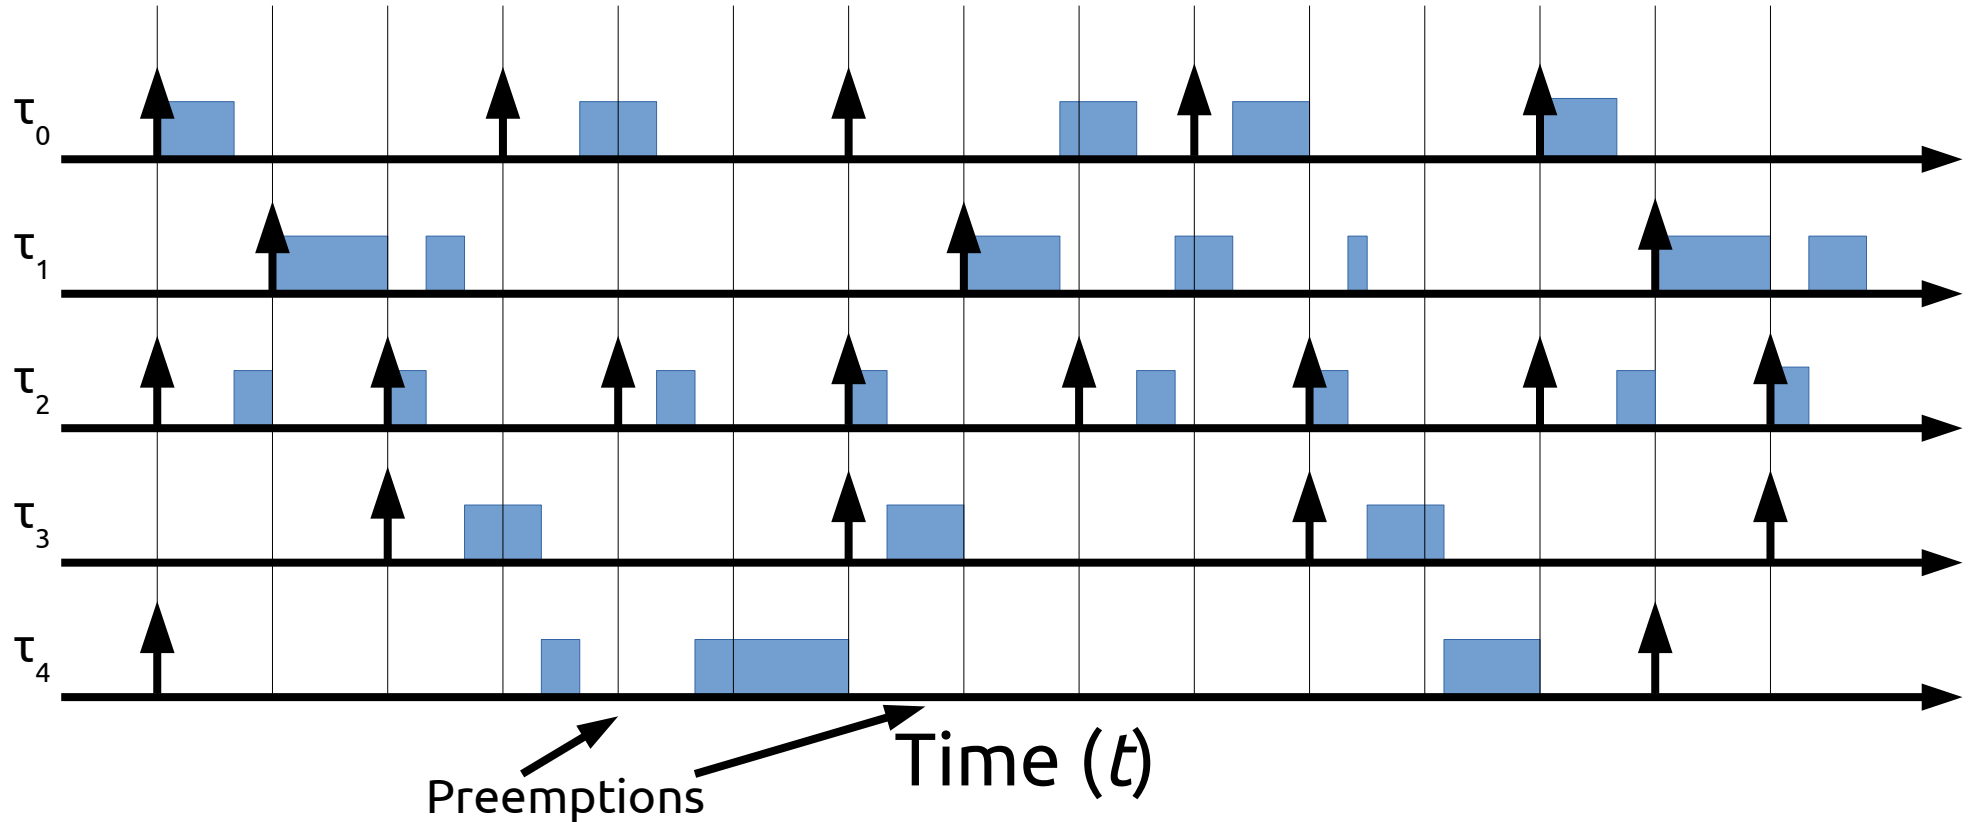
Task Model, Graphically



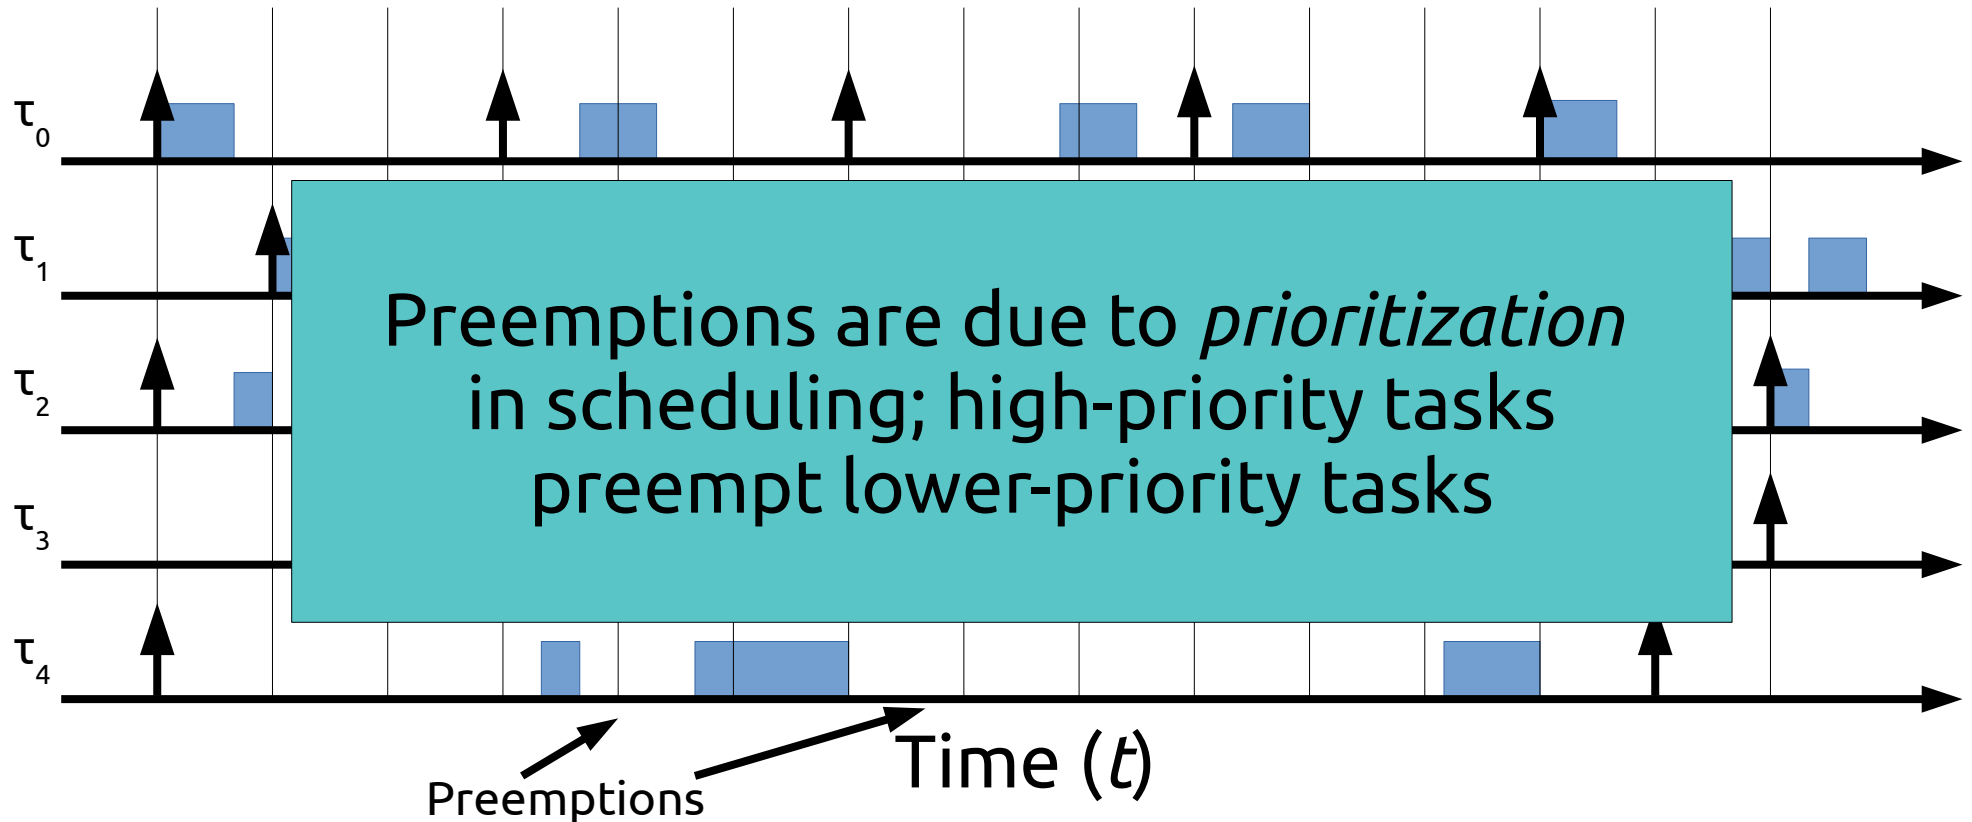
Task Model, Graphically



Task Model, Graphically



Task Model, Graphically



Priority

- Fixed priority scheduling (FP):
 - each task has a priority (σ_i)
 - highest priority, active task always executed
- Dynamic priority scheduling:
 - earliest deadline first (EDF)
 - always execute task with *earliest deadline*
- We'll only discuss (FP)

Fixed Priority Assignment

- How do we assign a priority, σ_i , to each task?
- Example:
 - τ_0 : $e_0 = 2$, $p_0 = 5$ with *low priority*
 - τ_1 : $e_1 = 4$, $p_1 = 10$ with *high priority*
- Note: only 8/10 time units spent on execution
- *Will we meet all deadlines?*

Fixed Priority Assignment


- How do we assign a priority, σ_i , to each task?
- *Ideas?*

Fixed Priority Assignment

- How do we assign a priority, σ_i , to each task?
- One option: “Rate-Monotonic” assignment
 - Lower p_i = higher priority
 - Faster rate = higher priority
 - Infrequent execution = lower priority

Why task models?

- A model lets you talk about a *specifically-specified idea* rather than *messy implementations*
- *Goal*: real-time systems
 - **Implementation** → policies to instantiate a model
 - **Model** → mathematical description of impl.
 - **Analysis** → derive properties from model



- *Goal*: real-time systems
 - **Implementation** → policies to instantiate a model
 - **Model** → mathematical description of impl.
 - **Analysis** → derive properties from model

Analysis: Schedulability

- Simple question:
Can the system meet every task's deadlines?
- A mathematical analysis
- ...using the properties of the task model
- ...that describes the implementation
- ...to make sure the RT system controls the world *predictably*

Analysis: Response-Time Analysis (RTA)

Assumes Fixed-Priority Scheduling

- What is the maximum bound to complete exec?
- Intuition: Over a window of time (w), calculate the *interference* ($if_i(w)$) from higher-prio tasks

$$if_i(p_i) + e_i > p_i$$

→ cannot meet deadline

→ system not schedulable

How would you calculate interference?

First scenario: Interference for τ_0

- τ_0 : $e_0 = 2$, $p_0 = 5$ with *low priority*
- τ_1 : $e_1 = 4$, $p_1 = 10$ with *high priority*

How would you calculate interference?

Second scenario: Interference for τ_2

- τ_0 : $e_0 = 1$, $p_0 = 3$ with *high priority*
- τ_1 : $e_1 = 3$, $p_1 = 6$ with *middle priority*
- τ_2 : $e_2 = 1$, $p_2 = 9$ with *low priority*

Response-Time Analysis

- Recall, we want to know if: $if_i(p_i) + e_i > p_i$
- $RTA_i(w)$ = execution time and interference

$$RTA_i(w) = e_i + \sum_{HP(i) \ni j} \lceil RTA_i(w)/p_j \rceil e_j$$

$HP(i) \rightarrow$ all tasks with higher priority than τ_i

Response-Time Analysis

$$RTA_i(w) = e_i + \sum_{HP(i) \ni j} \lceil RTA_i(w)/p_j \rceil e_j$$

HP(i) → all tasks with higher priority than τ_i

Algorithm:

$p = e_i$

$r = RTA_i(p)$

while ($p \neq r$):

$p = r$

$r = RTA_i(r)$

// start by seeing RTA over the exec time

// have we **converged**?

// compute the next span

Interference for τ_2

- τ_0 : $e_0 = 1$, $p_0 = 3$ with *high priority*
- τ_1 : $e_1 = 3$, $p_1 = 6$ with *middle prio*
- τ_2 : $e_2 = 1$, $p_2 = 9$ with *low priority*

$$RTA_i(w) = e_i + \sum_{HP(i) \ni j} \lceil RTA_i(w)/p_j \rceil e_j$$

$$RTA_2(1) = 1 + \lceil 1/6 \rceil 3 + \lceil 1/3 \rceil 1 = 1 + 3 + 1 = 5$$

$$RTA_2(5) = 1 + \lceil 5/6 \rceil 3 + \lceil 5/3 \rceil 1 = 1 + 3 + 2 = 6$$

$$RTA_2(6) = 1 + \lceil 6/6 \rceil 3 + \lceil 6/3 \rceil 1 = 1 + 3 + 2 = 6$$

Interference for τ_2

- τ_0 : $e_0 = 1$, $p_0 = 3$ with *high priority*
- τ_1 : $e_1 = 3$, $p_1 = 6$ with *middle prio*
- τ_2 : $e_2 = 1$, $p_2 = 9$ with *low priority*

$$RTA_i(w) = e_i + \sum_{HP(i) \ni j} \lceil RTA_i(w)/p_j \rceil e_j$$

$$RTA_2(1) = 1 + \lceil 1/6 \rceil 3 + \lceil 1/3 \rceil 1 = 1 + 3 + 1 = 5$$

$$RTA_2(5) = 1 + \lceil 5/6 \rceil 3 + \lceil 5/3 \rceil 1 = 1 + 3 + 2 = 6$$

$$RTA_2(\mathbf{6}) = 1 + \lceil 6/6 \rceil 3 + \lceil 6/3 \rceil 1 = 1 + 3 + 2 = \mathbf{6}$$

Interference for τ_2

- τ_0 : $e_0 = 1$, $p_0 = 3$ with *high priority*
- τ_1 : $e_1 = 3$, $p_1 = 6$ with *middle prio*
- τ_2 : $e_2 = 2$, $p_2 = 9$ with *low priority*

$$RTA_i(w) = e_i + \sum_{HP(i) \ni j} \lceil RTA_i(w)/p_j \rceil e_j$$

$$RTA_2(2) = 2 + \lceil 2/6 \rceil 3 + \lceil 2/3 \rceil 1 = 2 + 3 + 1 = 6$$

$$RTA_2(6) = 2 + \lceil 6/6 \rceil 3 + \lceil 6/3 \rceil 1 = 2 + 3 + 2 = 7$$

$$RTA_2(7) = 2 + \lceil 7/6 \rceil 3 + \lceil 7/3 \rceil 1 = 2 + 6 + 3 = 11$$

$$11 (RTA_2) > 9 (p_2)$$

Interference for τ_2

- τ_0 : $e_0 = 1$, $p_0 = 3$ with *high priority*
- τ_1 : $e_1 = 3$, $p_1 = 6$ with *middle prio*
- τ_2 : $e_2 = 2$, $p_2 = 9$ with *low priority*

$$RTA_i(w) = e_i + \sum_{HP(i) \ni j} \lceil RTA_i(w)/p_j \rceil e_j$$

$$RTA_2(2) = 2 + \lceil 2/6 \rceil 3 + \lceil 2/3 \rceil 1 = 2 + 3 + 1 = 6$$

$$RTA_2(6) = 2 + \lceil 6/6 \rceil 3 + \lceil 6/3 \rceil 1 = 2 + 3 + 2 = 7$$

$$RTA_2(7) = 2 + \lceil 7/6 \rceil 3 + \lceil 7/3 \rceil 1 = 2 + 6 + 3 = 11$$

$$11 (RTA_2) > 9 (p_2)$$

Backing up: *What have we done?*

- 1) Created a task model
- 2) Backed up by an implementation (FP sched)
- 3) Analyzed the task model
 - RTA: pass for all tasks → system is *schedulable*

Additional Concerns

- “Shared resources” – locks around shared data-structures between tasks
 - 1) Low-priority task takes lock
 - 2) Preemption, switch to high-priority task
 - 3) High-priority task attempts to take lock→ *high-prio task suffers low-prio interference*
- Task switches aren't free; how analyze?
→ *all tasks suffer from dispatch latency inteference*

Additional Concerns

- “Shared resources” – locks around shared data-structures between tasks*
- Task switches aren't free

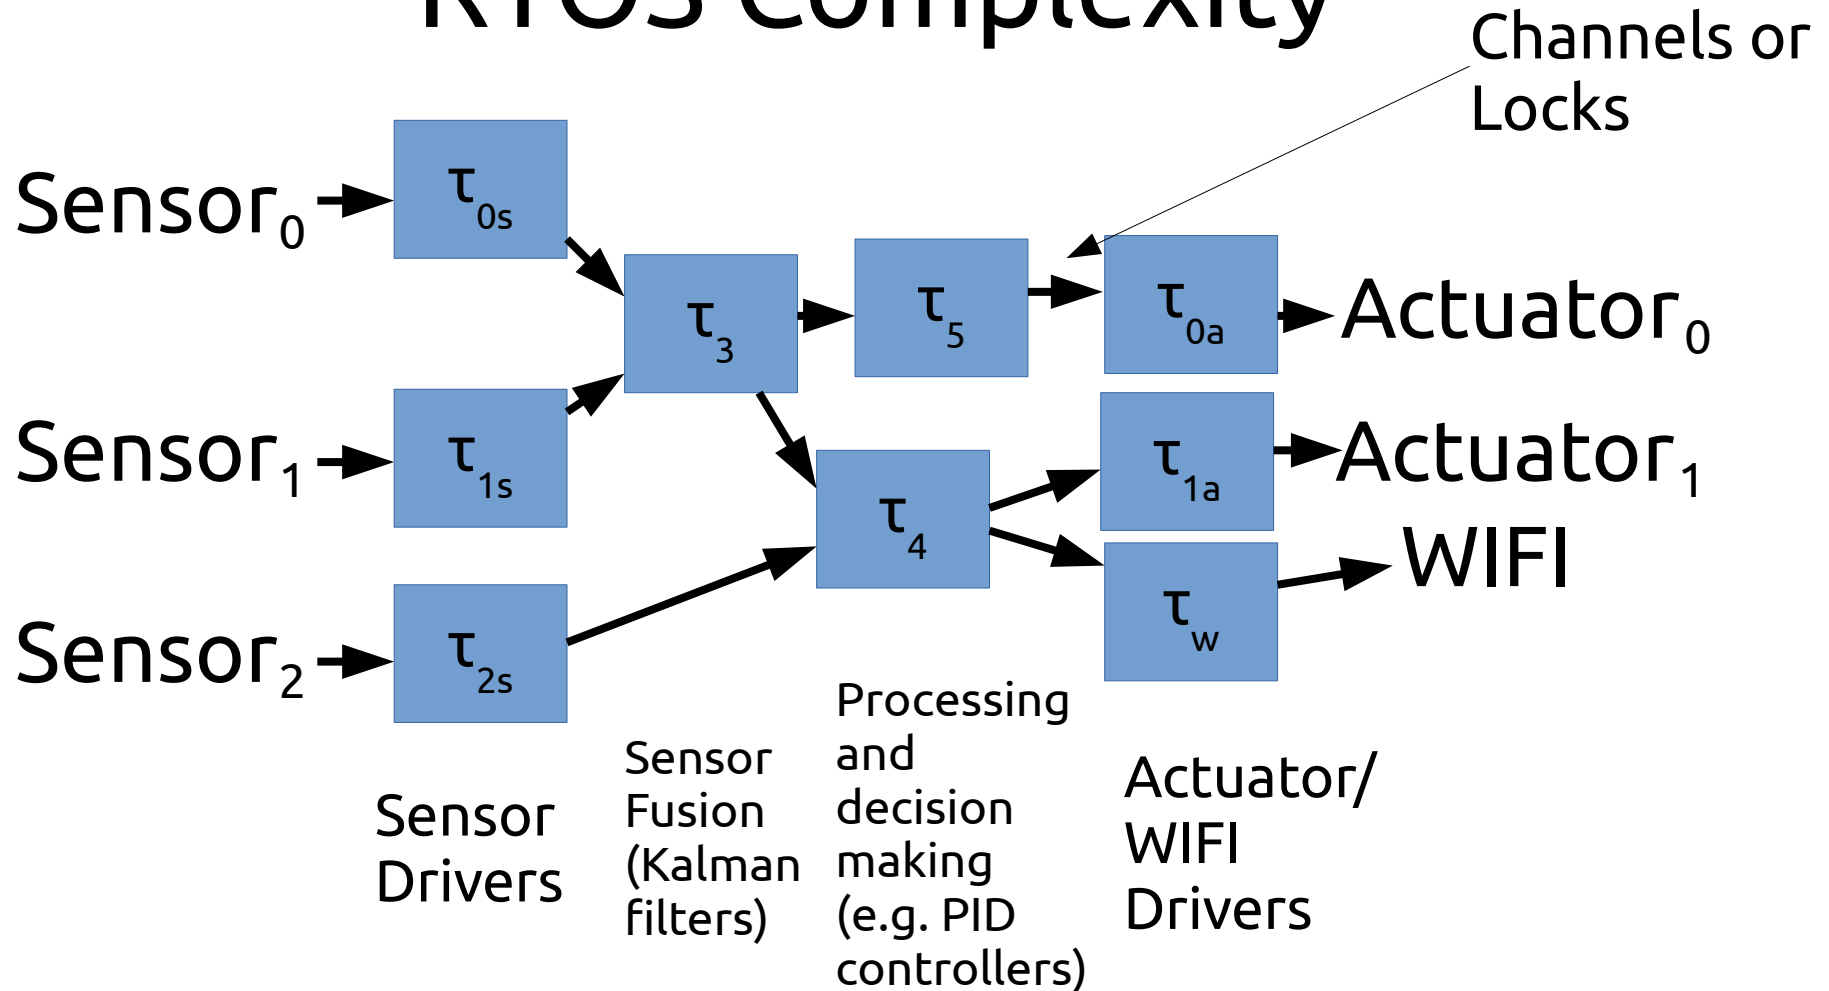
$$RTA_i(w) = e_i + e_{\text{lock}} + \sum_{HP(i) \ni j} \lceil RTA_i(w)/p_j \rceil (e_j + 2(e_{\text{swtch}}))$$

* Assumes predictable resource sharing protocol (e.g. Priority Inheritance)

Multicore

- Partitioned: Each core has its own set of tasks, and only schedules those
 - Just do an RTA per core!
- Global: All cores share the same runqueue of tasks
 - Not common
 - Need another analysis

RTOS Complexity



RTOS Complexity

Channels or
Locks

Sensor → τ_{os}

Mars Curiosity Rover: ~240 Tasks!

Sensor
Drivers

Sensor
Fusion
(Kalman
filters)

and
decision
making
(e.g. PID
controllers)

Actuator/
WIFI
Drivers